

# Recurrent Neural Networks

Deep Learning – Lecture 5

Efstratios Gavves



UNIVERSITY OF AMSTERDAM

# Sequential Data

---

So far, all tasks assumed ***stationary*** data



Neither all data, nor all tasks are stationary though

# Sequential Data: Text

---



What

# Sequential Data: Text

---



What about

# Sequential Data: Text

---



What about inputs that appear in sequences, such as text? Could a neural network handle such modalities?

# Memory needed

---



$$\Pr(x) = \prod_i \Pr(x_i | x_1, \dots, x_{i-1})$$

What about inputs that appear in sequences, such as text? Could a neural network handle such modalities?



# Sequential data: Video

---



# Quiz: Other sequential data?



# Quiz: Other sequential data?

## Time series data

- ☐ Stock exchange
- ☐ Biological measurements
- ☐ Climate measurements
- ☐ Market analysis

## Speech/Music

## User behavior in websites


...

# Applications

[Click to go to the video in Youtube](#)



NeuralTalk and Walk, recognition, text description of the image while walking



[Click to go to the website](#)

## CloudCV: Visual Question Answering (VQA)

More details about the VQA dataset can be found [here](#).  
State-of-the-art VQA model and code available [here](#).

CloudCV can answer questions you ask about an image.  
Browsers currently supported: Google Chrome, Mozilla Firefox.

### Try CloudCV VQA: Sample Images

Click on one of these images to send it to our servers (Or upload your own images below).



Hi Motherboard readers!  
This entire post was hand written by a neural network.  
(It probably writes better than you.)  
Of course, a neural network doesn't actually have hands.  
And the original text was typed by me, a human.  
So what's going on here?  
A neural network is a program that can learn to follow a set of rules.  
But it can't do it alone. It needs to be trained.  
This neural network was trained on a corpus of writing samples.

... samples of actual hand-writing,  
but of the locations of a pen-tip as people write.  
This is how the network learns and creates different styles  
from prior examples.  
And it can use the knowledge  
to generate handwritten notes from inputted text.  
It can create its own style, or mimic another's.  
No two notes are the same.  
It's the work of Alex Graves at the University of Toronto.  
And you can try it too!

# Machine Translation

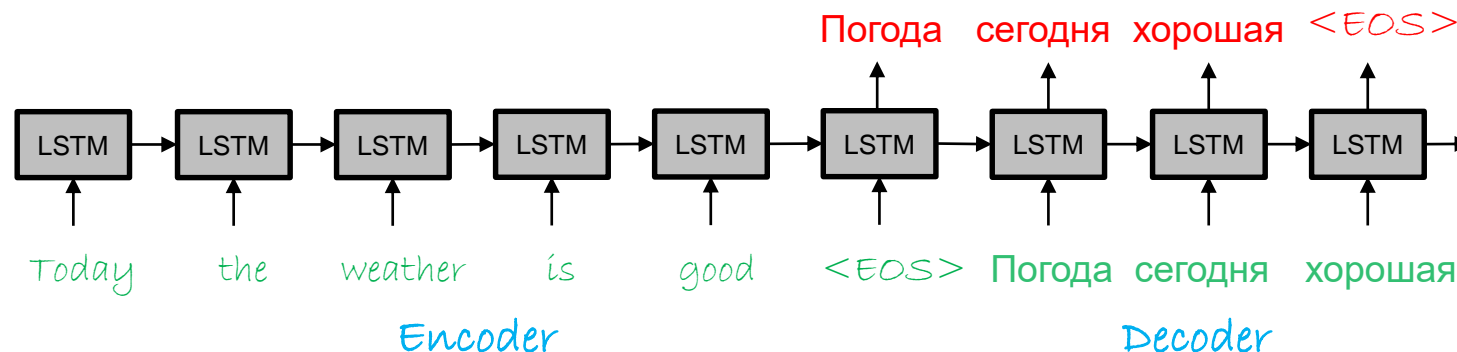
---

The phrase in the source language is one sequence

- “Today the weather is good”

The phrase in the target language is also a sequence

- “Погода сегодня хорошая”



# Image captioning

---

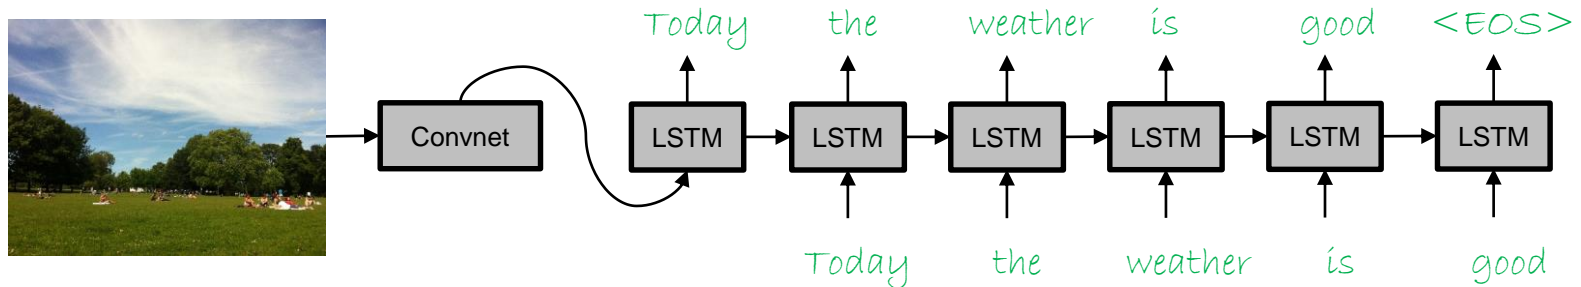
An image is a thousand words, literally!

Pretty much the same as machine translation

Replace the encoder part with the output of a Convnet

– E.g. use Alexnet or a VGG16 network

Keep the decoder part to operate as a translator



# Demo

---

[Click to go to the video in Youtube](#)



NeuralTalk and Walk, recognition, text description of the image while walking

# Question answering

---

Bleeding-edge research, no real consensus

- Very interesting open, research problems

Again, pretty much like machine translation

Again, Encoder-Decoder paradigm

- Insert the question to the encoder part
- Model the answer at the decoder part

Question answering with images also

- Again, bleeding-edge research
- How/where to add the image?
- What has been working so far is to add the image only in the beginning

Q: John entered the living room, where he met Mary. She was drinking some wine and watching a movie. What room did John enter?  
A: John entered the living room.



Q: what are the people playing?  
A: They play beach football



# Demo

---

[Click to go to the website](#)

## CloudCV: Visual Question Answering (VQA)

More details about the VQA dataset can be found [here](#).

State-of-the-art VQA model and code available [here](#)

CloudCV can answer questions you ask about an image

Browsers currently supported: Google Chrome, Mozilla Firefox

---

## Try CloudCV VQA: Sample Images

Click on one of these images to send it to our servers (Or [upload](#) your own images below)



# Handwriting

---

[Click to go to the website](#)

Hi Motherboard readers!

This entire post was hand written by a neural network.

(It probably writes better than you.)

Of course, a neural network doesn't actually have hands

And the original text was typed by me, a human.

So what's going on here?

A neural network is a program that can learn to follow a set of rules

But it can't do it alone. It needs to be trained.

This neural network was trained on a corpus of writing samples.

- a corpus is a set of actual hand-writing, of the locations of a pen-tip as people write.

how the network learns and creates different styles, from prior examples.

And it can use this knowledge to generate handwritten notes from inputted text.

~ create its own style, or mimic another's.


No two notes are the same.

the work of Alex Graves at the University of Toronto

And you can try it too!

# Recurrent Networks

---



# Sequences

---

Next data depend on previous data

Roughly equivalent to predicting what comes next

$$\Pr(x) = \prod_i \Pr(x_i | x_1, \dots, x_{i-1})$$



What about inputs that appear in  
sequences, such as text? Could a neural  
network handle such modalities?

# Why sequences?

---

Parameters are reused → Fewer parameters → Easier modelling

Generalizes well to arbitrary lengths → Not constrained to specific length

*RecurrentModel( I think, therefore, I am! )*

≡

*RecurrentModel( Everything is repeated, in a circle. History is a master because it teaches us that it doesn't exist. It's the permutations that matter.)*

However, often we pick a “frame”  $T$  instead of an arbitrary length

$$\Pr(x) = \prod_i \Pr(x_i | x_{i-T}, \dots, x_{i-1})$$

# Quiz: What is really a sequence?

Data inside a sequence are ... ?

I am Bond , James

McGuire

Bond

tired

am

!



# Quiz: What is really a sequence?

Data inside a sequence are non i.i.d.

- Identically, independently distributed

The next “word” depends on the previous “words”

- Ideally on all of them

We need **context**, and we need **memory**!

How to model context and memory ?

I am Bond , James

Bond

McGuire

Bond

tired

am

!

# $x_i \equiv$ One-hot vectors

---

A vector with all zeros except for the active dimension

12 words in a sequence  $\rightarrow$  12 One-hot vectors

After the one-hot vectors apply an embedding

□ Word2Vec, GloVE

<u>vocabulary</u>	<u>One-hot vectors</u>							
I	I	1	I	0	I	0	I	0
am	am	0	am	1	am	0	am	0
Bond	Bond	0	Bond	0	Bond	1	Bond	0
James	James	0	James	0	James	0	James	1
tired	tired	0	tired	0	tired	0	tired	0
,	,	0	,	0	,	0	,	0
McGuire	McGuire	0	McGuire	0	McGuire	0	McGuire	0
!	!	0	!	0	!	0	!	0

# Quiz: Why not just indices?

One-hot representation

OR?

Index representation

$$x_{t=1,2,3,4} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 \end{bmatrix}$$

1    am   James   McGuire

1    am   James   McGuire

$$x_{\text{"I"}} = 1$$

$$x_{\text{"am"}} = 2$$

$$x_{\text{"James"}} = 4$$

$$x_{\text{"McGuire"}} = 7$$

# Quiz: Why not just indices?

One-hot representation

OR?

Index representation

I	am	James	McGuire
1	0	0	0
0	1	0	0
0	0	0	0
0	0	1	0
0	0	0	0
0	0	0	0
0	0	0	1
0	0	0	0

$x_{\text{"I"}}$   $x_{\text{"James"}}$   
 $x_{\text{"am"}}$   $x_{\text{"McGuire"}}$

$$\text{distance}(x_{\text{"am"}}, x_{\text{"McGuire"}}) = 1$$

$$\text{distance}(x_{\text{"I"}}, x_{\text{"am"}}) = 1$$

I	am	James	McGuire
---	----	-------	---------

$$q_{\text{"I"}} = 1$$

$$q_{\text{"am"}} = 2$$

$$q_{\text{"James"}} = 4$$

$$q_{\text{"McGuire"}} = 7$$

$$\text{distance}(q_{\text{"am"}}, q_{\text{"McGuire"}}) = 5$$

$$\text{distance}(q_{\text{"I"}}, q_{\text{"am"}}) = 1$$

No, because then some words get closer together for no good reason (artificial bias between words)

# Memory

---

A representation of the past

A memory must project information at timestep  $t$  on a latent space  $c_t$  using parameters  $\theta$

Then, re-use the projected information from  $t$  at  $t + 1$

$$c_{t+1} = h(x_{t+1}, c_t; \theta)$$

Memory parameters  $\theta$  are shared for all timesteps  $t = 0, \dots$


$$c_{t+1} = h(x_{t+1}, h(x_t, h(x_{t-1}, \dots h(x_1, c_0; \theta); \theta); \theta); \theta))$$

# Memory as a Graph

---

## Simplest model

- Input with parameters  $U$
- Memory embedding with parameters  $W$
- Output with parameters  $V$






# Memory as a Graph

---

## Simplest model

- Input with parameters  $U$
- Memory embedding with parameters  $W$
- Output with parameters  $V$




# Memory as a Graph

---

## Simplest model


- Input with parameters  $U$
- Memory embedding with parameters  $W$
- Output with parameters  $V$




# Folding the memory

---

unrolled/unfolded Network



Folded Network




# Recurrent Neural Network (RNN)

---

Only **two** equations

$$c_t = \tanh(U x_t + W c_{t-1})$$
$$y_t = \text{softmax}(V c_t)$$



# RNN Example

---

Vocabulary of 5 words

A memory of 3 units

- Hyperparameter that we choose like layer size
- $c_t: [3 \times 1]$ ,  $W: [3 \times 3]$

An input projection of 3 dimensions

- $U: [3 \times 5]$

An output projections of 10 dimensions

- $V: [10 \times 3]$

$$\textcolor{red}{U} \cdot x_{t=4} = \begin{bmatrix} 0.1 & -0.3 & 1.2 & 0.6 & -0.8 \\ -0.2 & 0.4 & 0.5 & 0.9 & -0.1 \\ -0.1 & 0.2 & -0.7 & -0.8 & 0.3 \end{bmatrix} \cdot \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \\ 0 \end{bmatrix} = \begin{bmatrix} 0.6 \\ 0.9 \\ -0.8 \end{bmatrix} = U^{(4)}$$



$$c_t = \tanh(\textcolor{red}{U} x_t + \textcolor{blue}{W} c_{t-1})$$

$$y_t = \text{softmax}(\textcolor{green}{V} c_t)$$

# Rolled Network vs. MLP?

What is really different?

- Steps instead of layers
- Step parameters shared in Recurrent Network
- In a Multi-Layer Network parameters are different







# Rolled Network vs. MLP?

What is really different?

- Steps instead of layers
- Step parameters shared in Recurrent Network
- In a Multi-Layer Network parameters are different





# Rolled Network vs. MLP?

What is really different?

- Steps instead of layers
- Step parameters shared in Recurrent Network
- In a Multi-Layer Network parameters are different



- Sometimes intermediate outputs are not even needed
- Removing them, we almost end up to a standard Neural Network



# Rolled Network vs. MLP?

What is really different?

- Steps instead of layers
- Step parameters shared in Recurrent Network
- In a Multi-Layer Network parameters are different



# Training Recurrent Networks

---

Cross-entropy loss

$$P = \prod_{t,k} y_{tk}^{l_{tk}} \Rightarrow \mathcal{L} = -\log P = \sum_t \mathcal{L}_t = -\frac{1}{T} \sum_t l_t \log y_t$$

Backpropagation Through Time (BPTT)

- Again, chain rule
- Only difference: Gradients survive over time steps

# Training RNNs is hard

---

Vanishing gradients

- After a few time steps the gradients become almost 0

Exploding gradients

- After a few time steps the gradients become huge

Can't capture long-term dependencies

# Alternative formulation for RNNs

---

An alternative formulation to derive conclusions and intuitions

$$c_t = W \cdot \tanh(c_{t-1}) + U \cdot x_t + b$$

$$\mathcal{L} = \sum_t \mathcal{L}_t(c_t)$$

# Another look at the gradients

---

$$\mathcal{L} = L(c_T(c_{T-1}(\dots(c_1(x_1, c_0; W); W); W); W); W)$$

$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_t}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W}$$

$$\frac{\partial \mathcal{L}}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} = \underbrace{\frac{\partial \mathcal{L}}{\partial c_t} \cdot \frac{\partial c_t}{\partial c_{t-1}} \cdot \frac{\partial c_{t-1}}{\partial c_{t-2}} \cdot \dots \cdot \frac{\partial c_{\tau+1}}{\partial c_\tau}}_{\text{Rest} \rightarrow \text{short-term factors}} \leq \eta^{t-\tau} \frac{\partial \mathcal{L}_t}{\partial c_t}$$

*Rest*  $\rightarrow$  short-term factors

$t \gg \tau \rightarrow$  long-term factors

The RNN gradient is a recursive product of  $\frac{\partial c_t}{\partial c_{t-1}}$

# Exploding/Vanishing gradients

---

$$\frac{\partial \mathcal{L}}{\partial c_t} = \frac{\partial \mathcal{L}}{\partial c_T} \cdot \underbrace{\frac{\partial c_T}{\partial c_{T-1}} \cdot \frac{\partial c_{T-1}}{\partial c_{T-2}} \cdot \dots \cdot \frac{\partial c_{t+1}}{\partial c_t}}_{\substack{< 1 \quad < 1 \quad < 1}} \left\{ \frac{\partial \mathcal{L}}{\partial w} \ll 1 \Rightarrow \text{Vanishing gradient} \right.$$

$$\frac{\partial \mathcal{L}}{\partial c_t} = \frac{\partial \mathcal{L}}{\partial c_T} \cdot \underbrace{\frac{\partial c_T}{\partial c_{T-1}} \cdot \frac{\partial c_{T-1}}{\partial c_{T-2}} \cdot \dots \cdot \frac{\partial c_1}{\partial c_t}}_{\substack{> 1 \quad > 1 \quad > 1}} \left\{ \frac{\partial \mathcal{L}}{\partial w} \gg 1 \Rightarrow \text{Exploding gradient} \right.$$



# Vanishing gradients

---

The gradient of the error w.r.t. to intermediate cell

$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_r}{\partial y_t} \frac{\partial y_t}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W}$$

$$\frac{\partial c_t}{\partial c_\tau} = \prod_{t \geq k \geq \tau} \frac{\partial c_k}{\partial c_{k-1}} = \prod_{t \geq k \geq \tau} W \cdot \partial \tanh(c_{k-1})$$

# Vanishing gradients

---

The gradient of the error w.r.t. to intermediate cell

$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_r}{\partial y_t} \frac{\partial y_t}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W}$$

$$\frac{\partial c_t}{\partial c_\tau} = \prod_{t \geq k \geq \tau} \frac{\partial c_k}{\partial c_{k-1}} = \prod_{t \geq k \geq \tau} W \cdot \partial \tanh(c_{k-1})$$

Long-term dependencies get exponentially smaller weights


# Gradient clipping for exploding gradients

---

Scale the gradients to a threshold

## Pseudocode

```
1.  $g \leftarrow \frac{\partial \mathcal{L}}{\partial W}$   
2. if  $\|g\| > \theta_0$ :  
     $g \leftarrow \frac{\theta_0}{\|g\|} g$   
    else:  
        print('Do nothing')
```



Simple, but works!

# Rescaling vanishing gradients?

---

Not good solution

Weights are shared between timesteps  $\rightarrow$  Loss summed over timesteps


$$\mathcal{L} = \sum_t \mathcal{L}_t \Rightarrow \frac{\partial \mathcal{L}}{\partial W} = \sum_t \frac{\partial \mathcal{L}_t}{\partial W}$$
$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_t}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W}$$

Rescaling for one timestep ( $\frac{\partial \mathcal{L}_t}{\partial W}$ ) affects all timesteps

- The rescaling factor for one timestep does not work for another

# More intuitively

---



# More intuitively

Let's say  $\frac{\partial \mathcal{L}_1}{\partial W} \propto 1, \frac{\partial \mathcal{L}_2}{\partial W} \propto 1/10, \frac{\partial \mathcal{L}_3}{\partial W} \propto 1/100, \frac{\partial \mathcal{L}_4}{\partial W} \propto 1/1000, \frac{\partial \mathcal{L}_5}{\partial W} \propto 1/10000$


$$\frac{\partial \mathcal{L}}{\partial W} = \sum_r \frac{\partial \mathcal{L}_r}{\partial W} = 1.1111$$

If  $\frac{\partial \mathcal{L}}{\partial W}$  rescaled to 1  $\rightarrow \frac{\partial \mathcal{L}_5}{\partial W} \propto 10^{-5}$

Longer-term dependencies negligible

- Weak recurrent modelling
- Learning focuses on the short-term only

$$\frac{\partial \mathcal{L}}{\partial W} = \frac{\partial \mathcal{L}_1}{\partial W} + \frac{\partial \mathcal{L}_2}{\partial W} + \frac{\partial \mathcal{L}_3}{\partial W} + \frac{\partial \mathcal{L}_4}{\partial W} + \frac{\partial \mathcal{L}_5}{\partial W}$$



# Recurrent networks $\propto$ Chaotic systems




Figure 4. This diagram illustrates how the change in  $\mathbf{x}_t$ ,  $\Delta \mathbf{x}_t$ , can be large for a small  $\Delta \mathbf{x}_0$ . The blue vs red (left vs right) trajectories are generated by the same maps  $F_1, F_2, \dots$  for two different initial states.




Figure 5. Illustrates how one can break apart the maps  $F_1, \dots, F_t$  into a constant map  $\tilde{F}$  and the maps  $U_1, \dots, U_t$ . The dotted vertical line represents the boundary between basins of attraction, and the straight dashed arrow the direction of the map  $\tilde{F}$  on each side of the boundary. This diagram is an extension of Fig. 4.

# Fixing vanishing gradients

---

Regularization on the recurrent weights

- Force the error signal not to vanish

$$\Omega = \sum_t \Omega_t = \sum_t \left( \frac{\left| \frac{\partial \mathcal{L}}{\partial c_{t+1}} \frac{\partial c_{t+1}}{\partial c_t} \right|}{\left| \frac{\partial \mathcal{L}}{\partial c_{t+1}} \right|} - 1 \right)^2$$

Advanced recurrent modules

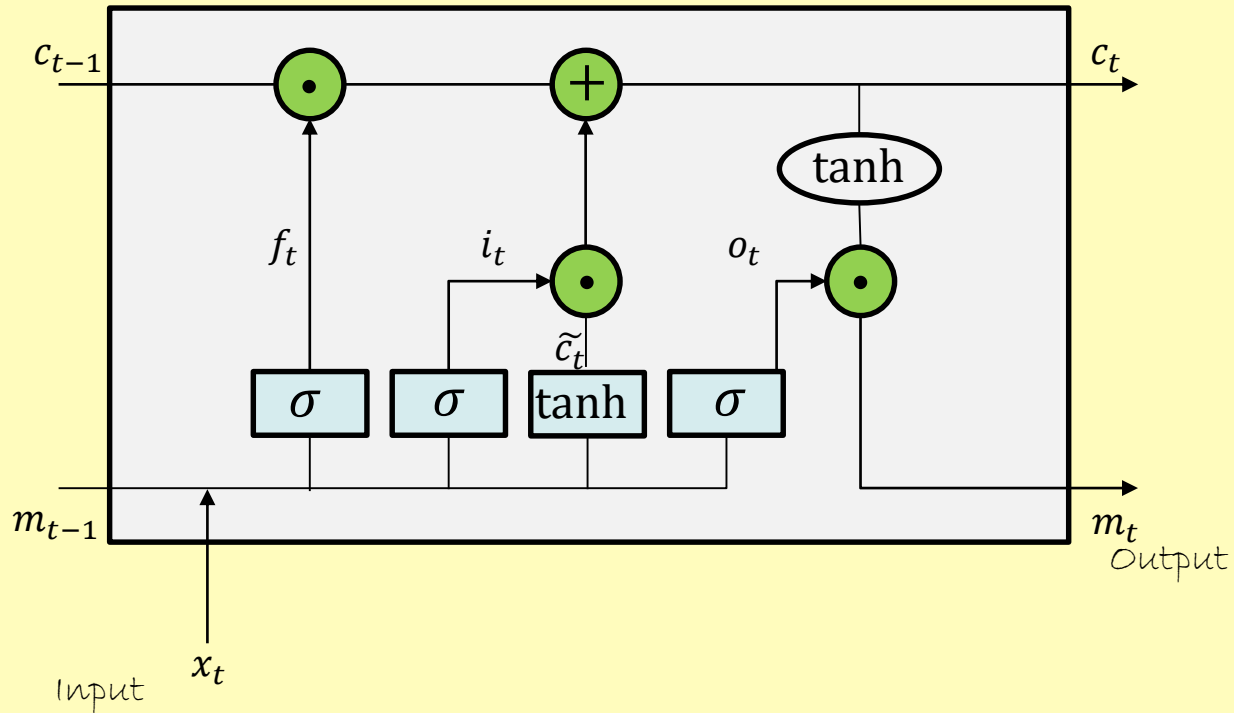
Long-Short Term Memory module

Gated Recurrent Unit module



# Advanced Recurrent Nets

---



# How to fix the vanishing gradients?

---

Error signal over time must have not too large, not too small norm

Let's have a look at the loss function

$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_r}{\partial y_t} \frac{\partial y_t}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W}$$
$$\frac{\partial c_t}{\partial c_\tau} = \prod_{t \geq k \geq \tau} \frac{\partial c_k}{\partial c_{k-1}}$$

# How to fix the vanishing gradients?

Error signal over time must have not too large, not too small norm

Let's have a look at the loss function

$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_r}{\partial y_t} \frac{\partial y_t}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W}$$
$$\frac{\partial c_t}{\partial c_\tau} = \prod_{t \geq k \geq \tau} \frac{\partial c_k}{\partial c_{k-1}}$$

# How to fix the vanishing gradients?

---

Error signal over time must have not too large, not too small norm

Solution: have an activation function with gradient equal to 1

$$\frac{\partial \mathcal{L}_t}{\partial W} = \sum_{\tau=1}^t \frac{\partial \mathcal{L}_r}{\partial y_t} \frac{\partial y_t}{\partial c_t} \frac{\partial c_t}{\partial c_\tau} \frac{\partial c_\tau}{\partial W}$$
$$\frac{\partial c_t}{\partial c_\tau} = \prod_{t \geq k \geq \tau} \frac{\partial c_k}{\partial c_{k-1}}$$

- Identify function has a gradient equal to 1

By doing so, gradients do not become too small not too large

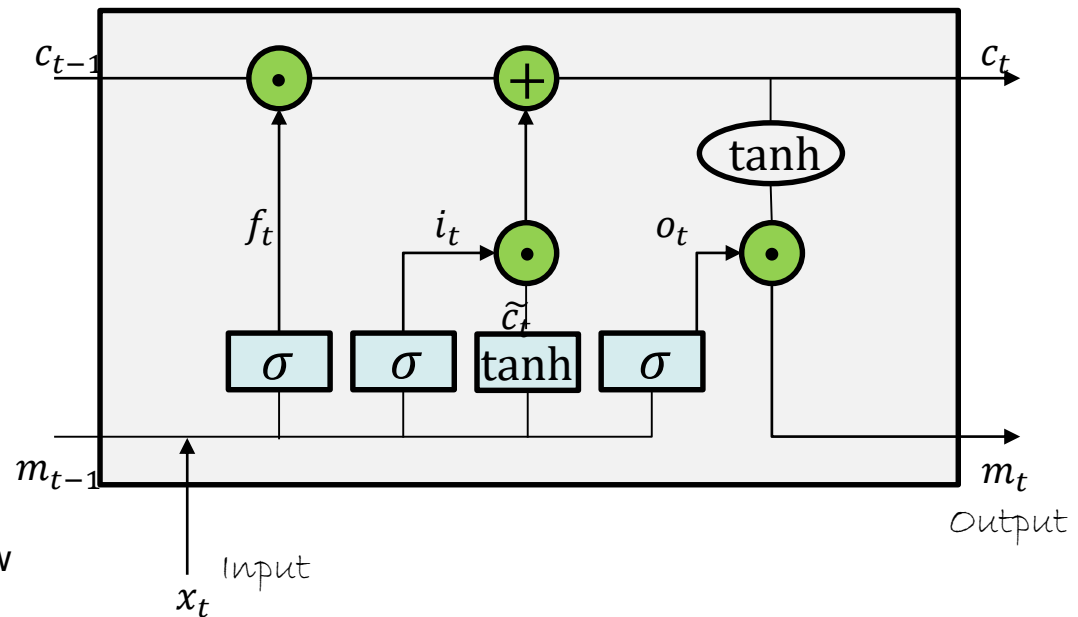
# Long Short-Term Memory

## LSTM: Beefed up RNN

Simple RNN {  $c_t = W \cdot \tanh(c_{t-1}) + U \cdot x_t + b$

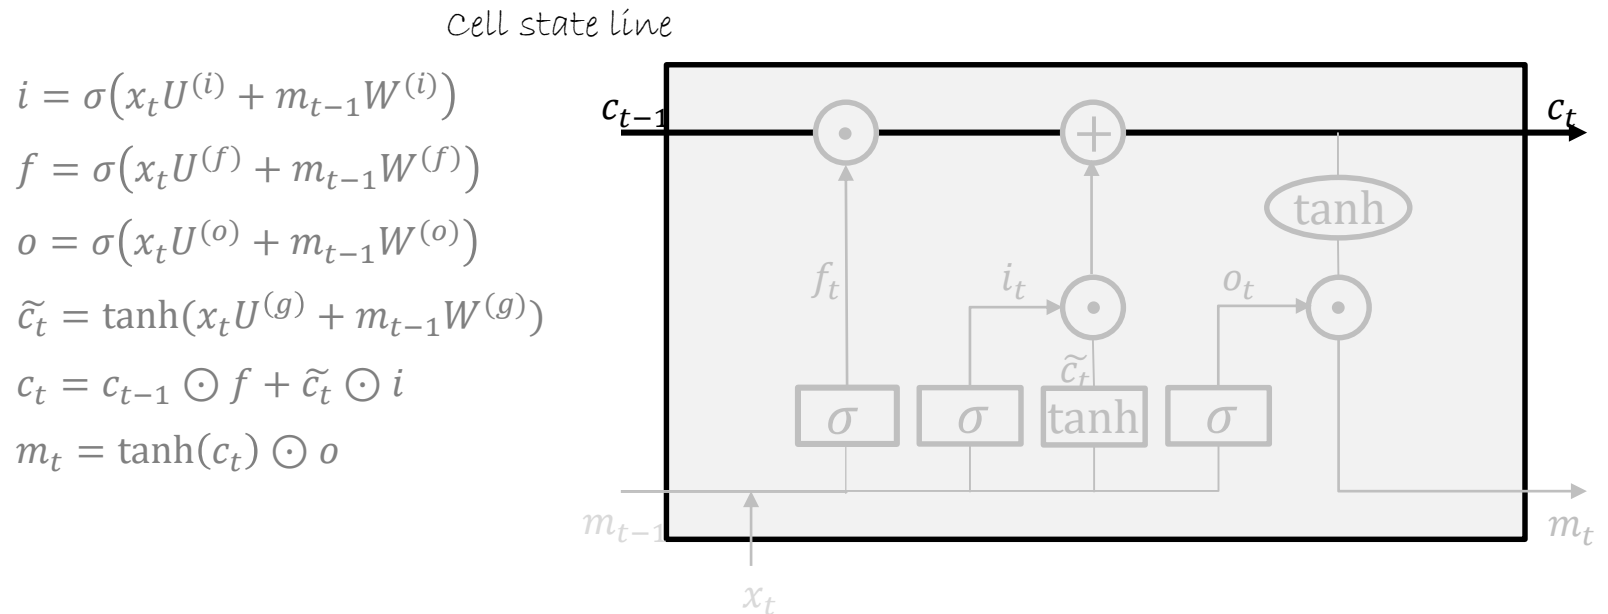
LSTM { 
$$\begin{aligned} i &= \sigma(x_t U^{(i)} + m_{t-1} W^{(i)}) \\ f &= \sigma(x_t U^{(f)} + m_{t-1} W^{(f)}) \\ o &= \sigma(x_t U^{(o)} + m_{t-1} W^{(o)}) \\ \tilde{c}_t &= \tanh(x_t U^{(g)} + m_{t-1} W^{(g)}) \\ c_t &= c_{t-1} \odot f + \tilde{c}_t \odot i \\ m_t &= \tanh(c_t) \odot o \end{aligned}$$

- The previous state  $c_{t-1}$  is connected to new  $c_t$  with no nonlinearity (identity function).
- The only other factor is the forget gate  $f$  which rescales the previous LSTM state.



# Cell state

The cell state carries the essential information over time



# LSTM nonlinearities

$\sigma \in (0, 1)$ : control gate – something like a switch

$\tanh \in (-1, 1)$ : recurrent nonlinearity

$$i = \sigma(x_t U^{(i)} + m_{t-1} W^{(i)})$$

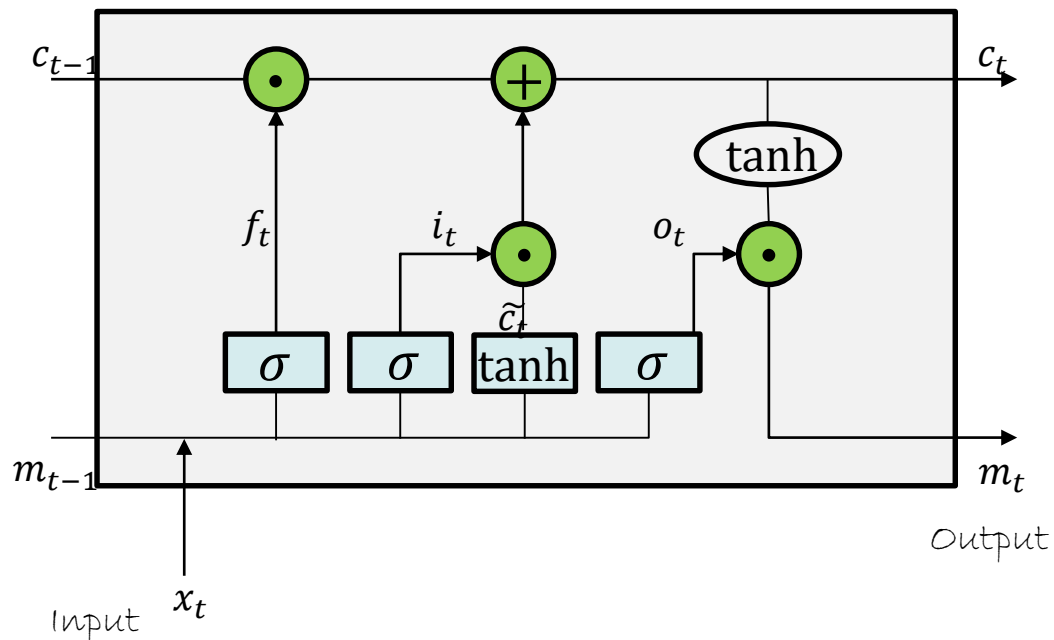
$$f = \sigma(x_t U^{(f)} + m_{t-1} W^{(f)})$$

$$o = \sigma(x_t U^{(o)} + m_{t-1} W^{(o)})$$

$$\tilde{c}_t = \tanh(x_t U^{(g)} + m_{t-1} W^{(g)})$$

$$c_t = c_{t-1} \odot f + \tilde{c}_t \odot i$$

$$m_t = \tanh(c_t) \odot o$$

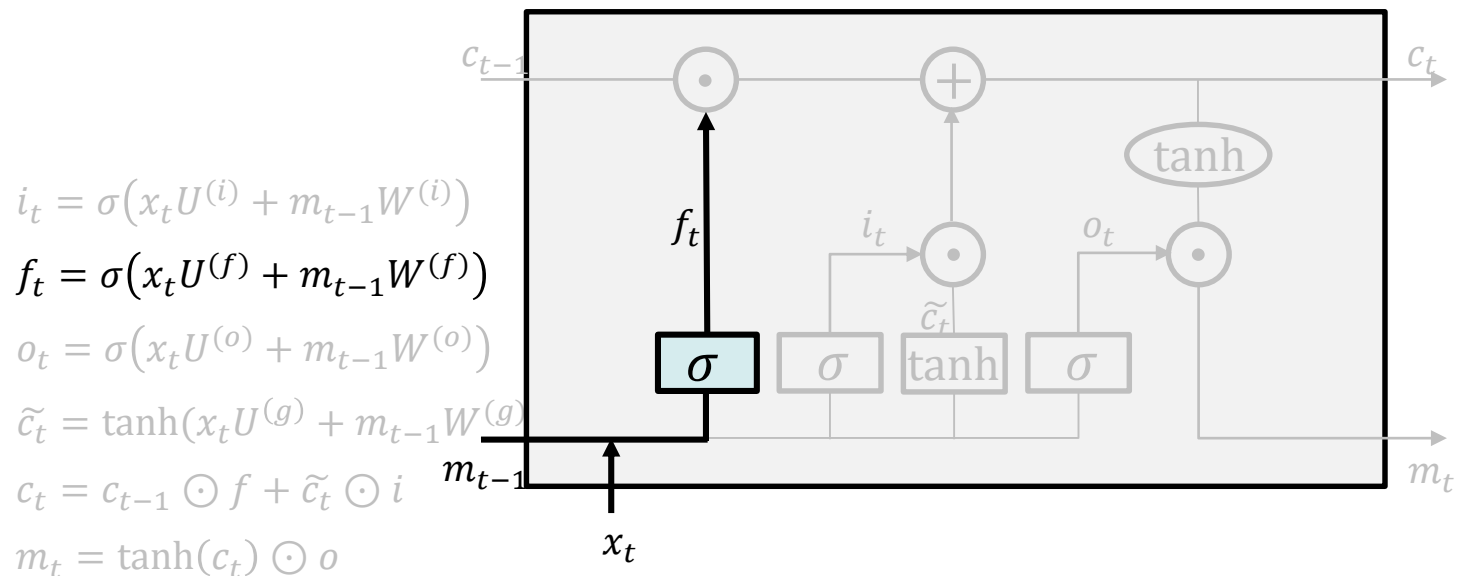


# LSTM Step-by-Step: Step (1)

E.g. LSTM on “Yesterday she slapped me. Today she loves me.”

Decide what to forget and what to remember for the new memory

- Sigmoid 1 → Remember everything
- Sigmoid 0 → Forget everything





# LSTM Step-by-Step: Step (2)

Decide what new information is relevant from the new input and should be added to the new memory

- Modulate the input  $i_t$
- Generate candidate memories  $\tilde{c}_t$

$$i_t = \sigma(x_t U^{(i)} + m_{t-1} W^{(i)})$$

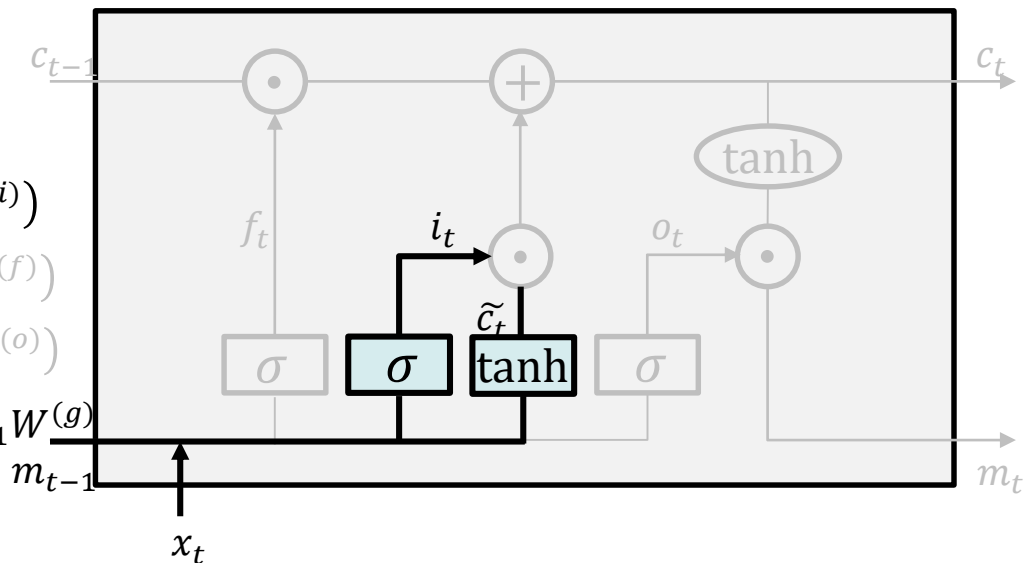
$$f_t = \sigma(x_t U^{(f)} + m_{t-1} W^{(f)})$$

$$o_t = \sigma(x_t U^{(o)} + m_{t-1} W^{(o)})$$

$$\tilde{c}_t = \tanh(x_t U^{(g)} + m_{t-1} W^{(g)})$$

$$c_t = c_{t-1} \odot f + \tilde{c}_t \odot i$$

$$m_t = \tanh(c_t) \odot o$$

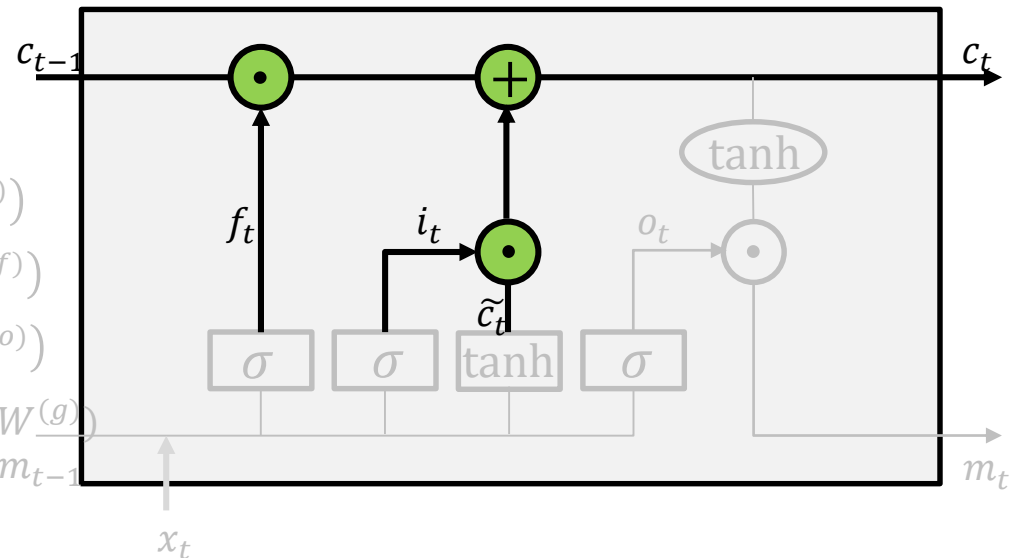


# LSTM Step-by-Step: Step (3)

Compute and update the current cell state  $c_t$

- Depends on the previous cell state
- What we decide to forget
- What inputs we allow
- The candidate memories

$$\begin{aligned}i_t &= \sigma(x_t U^{(i)} + m_{t-1} W^{(i)}) \\f_t &= \sigma(x_t U^{(f)} + m_{t-1} W^{(f)}) \\o_t &= \sigma(x_t U^{(o)} + m_{t-1} W^{(o)}) \\\tilde{c}_t &= \tanh(x_t U^{(g)} + m_{t-1} W^{(g)}) \\c_t &= c_{t-1} \odot f + \tilde{c}_t \odot i \\m_t &= \tanh(c_t) \odot o\end{aligned}$$



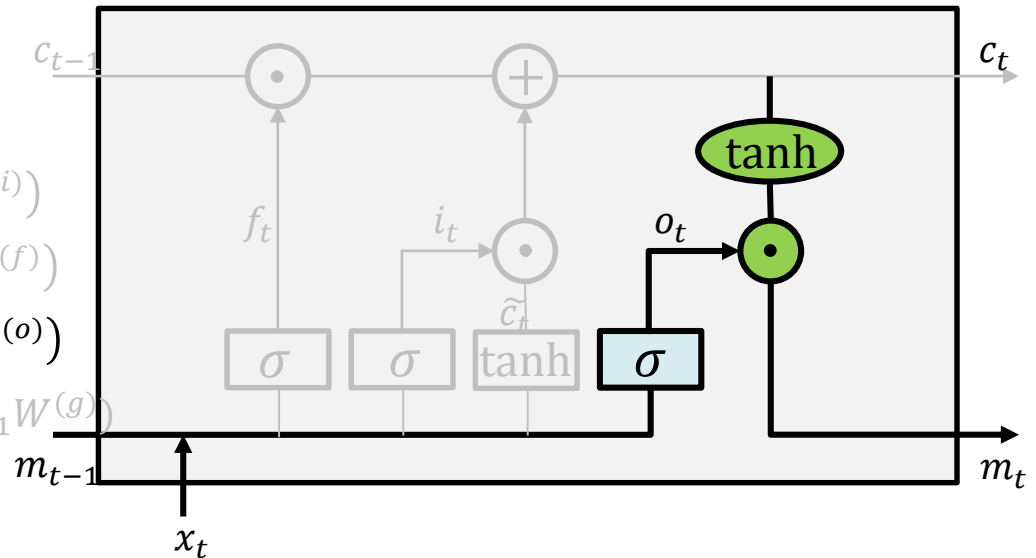
# LSTM Step-by-Step: Step (4)

Modulate the output

- Does the new cell state relevant?  $\rightarrow$  Sigmoid 1
- If not  $\rightarrow$  Sigmoid 0

Generate the new memory

$$\begin{aligned} i_t &= \sigma(x_t U^{(i)} + m_{t-1} W^{(i)}) \\ f_t &= \sigma(x_t U^{(f)} + m_{t-1} W^{(f)}) \\ o_t &= \sigma(x_t U^{(o)} + m_{t-1} W^{(o)}) \\ \tilde{c}_t &= \tanh(x_t U^{(g)} + m_{t-1} W^{(g)}) \\ c_t &= c_{t-1} \odot f + \tilde{c}_t \odot i \\ m_t &= \tanh(c_t) \odot o \end{aligned}$$



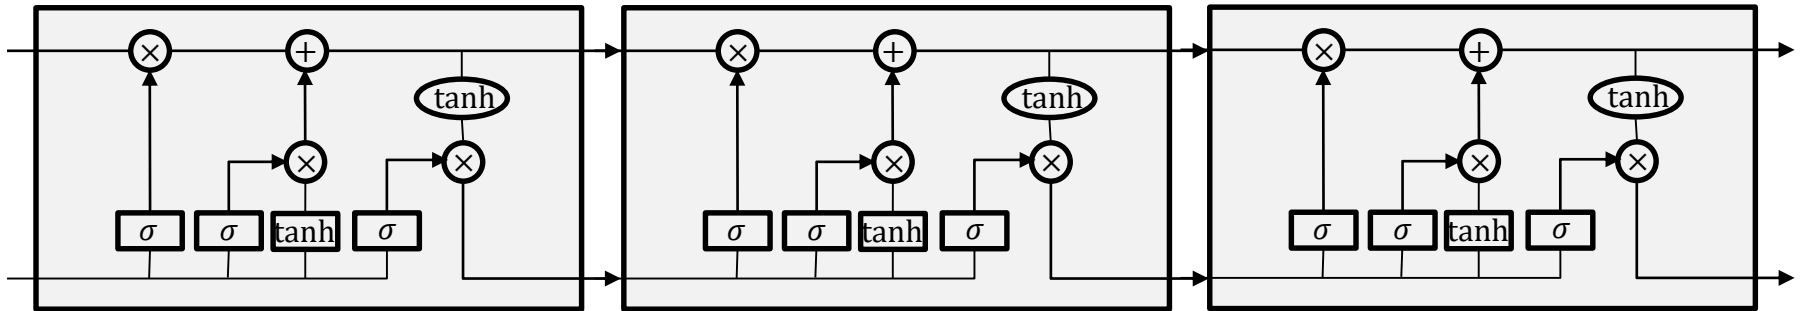
# LSTM Unrolled Network

---

Macroscopically very similar to standard RNNs

The engine is a bit different (more complicated)

- Because of their gates LSTMs capture long and short term dependencies



# Beyond RNN & LSTM


---

## LSTM with peephole connections

- Gates have access also to the previous cell states  $c_{t-1}$  (not only memories)
- Coupled forget and input gates,  $c_t = f_t \odot c_{t-1} + (1 - f_t) \odot \tilde{c}_t$
- Bi-directional recurrent networks

## Gated Recurrent Units (GRU)

## Deep recurrent architectures



## Recursive neural networks

- Tree structured

## Multiplicative interactions

## Generative recurrent architectures

# Take-away message

---

Recurrent Neural Networks (RNN) for sequences

Backpropagation Through Time

Vanishing and Exploding Gradients and Remedies

RNNs using Long Short-Term Memory (LSTM)

Applications of Recurrent Neural Networks

# Thank you!

---

